

Y86-64 ISA

'20H2

송 인 식

Outline

- Y86-64 ISA
 - Programmer-visible state
 - Instruction Format
 - Exceptions
- Y86-64 Programs
- CISC vs. RISC

Y86-64

- A strip-down version of x86-64 created by textbook authors for educational purposes
 - Similar state and instructions
 - Simpler encodings
 - Somewhere between CISC and RISC
- We will work through a CPU design example with Y86-64 in Chap. 4
- Y86-64 toolset available at <http://csapp.cs.cmu.edu/3e/sim.tar>
 - yas (assembler), yis (ISA simulator)
 - hcl2c (HCL to C translator), hcl2v (HCL to Verilog translator)
 - ssim, ssim+, psim (hardware simulator)

Y86-64 Programmer-Visible State

RF: Program registers

%rax	%rsp	%r8	%r12
%rcx	%rbp	%r9	%r13
%rdx	%rsi	%r10	%r14
%rbx	%rdi	%r11	

CC: Condition
codes

ZF	SF	OF
----	----	----

PC



Stat: Program status



DMEM: Memory



Y86-64 Programmer-Visible State

- Program Registers
 - 15 registers (no %r15), each 64 bits
- Condition codes
 - Single-bit flags set by arithmetic or logical instructions
 - ZF: Zero, SF: Negative, OF: Overflow
- Program Counter: address of next instruction
- Program Status: normal vs. error condition
- Memory
 - Byte-addressable storage array
 - Words stored in little-endian byte order

Y86-64 Instruction Set #1

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmovXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB	V					
rmmovq rA, D(rB)	4	0	rA	rB	D					
mrmmovq D(rB), rA	5	0	rA	rB	D					
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn	Dest							
call Dest	8	0	Dest							
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

Y86-64 Instructions

- 1 – 10 bytes of information read from memory
 - Can determine instruction length from first byte
- Only supports 64-bit operations
- RISC style
 - Not as many instruction types, and simpler encoding than with x86-64
 - Simple addressing mode: D(rA)
 - ALU instructions operate on registers (not memory)
 - Registers are specified in the fixed location, if any
- Each accesses and modifies some part(s) of the program state

Y86-64 [Conditional] Move Instructions

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmovXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB						
rmmovq rA, D(rB)	4	0	rA	rB						
rmrmovq D(rB), rA	5	0	rA	rB						
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn								
call Dest	8	0								
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

rrmovq	2	0
cmovle	2	1
cmovl	2	2
cmove	2	3
cmovne	2	4
cmovge	2	5
cmovg	2	6

Y86-64 ALU Instructions

Byte	0	1	2	3	4	5	6	7	8	9				
halt	0	0												
nop	1	0												
cmovXX rA, rB	2	fn	rA	rB										
irmovq V, rB	3	0	F	rB	V									
rmmovq rA, D(rB)	4	0	rA	rB	D									
mrmmovq D(rB), rA	5	0	rA	rB	D									
OPq rA, rB	6	fn	rA	rB	<div>addq<div>60</div></div> <div>subq<div>61</div></div> <div>andq<div>62</div></div> <div>xorq<div>63</div></div>									
jXX Dest	7	fn	Dest											
call Dest	8	0	Dest											
ret	9	0												
pushq rA	A	0	rA	F										
popq rA	B	0	rA	F										

Y86-64 [Conditional] Branch Instructions

Byte	0	1	2	3	4	5	6	7	8	9
halt	0	0								
nop	1	0								
cmovXX rA, rB	2	fn	rA	rB						
irmovq V, rB	3	0	F	rB	V					
rmmovq rA, D(rB)	4	0	rA	rB	D					
mrmmovq D(rB), rA	5	0	rA	rB	D					
OPq rA, rB	6	fn	rA	rB						
jXX Dest	7	fn	Dest							
call Dest	8	0	Dest							
ret	9	0								
pushq rA	A	0	rA	F						
popq rA	B	0	rA	F						

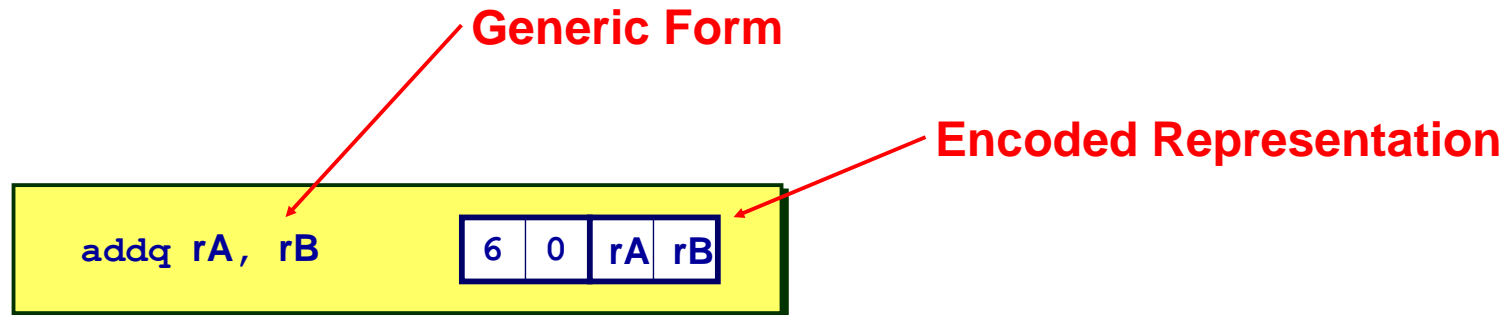
jmp	7	0
jle	7	1
j1	7	2
je	7	3
jne	7	4
jge	7	5
jg	7	6

Encoding Registers

- Each register has 4-bit ID
 - Same encoding as in x86-64
- Register ID 15 (0xf) indicates “no register”
 - Will use this in our hardware design in multiple places

Number	Register name	Number	Register name
0	%rax	8	%r8
1	%rcx	9	%r9
2	%rdx	A	%r10
3	%rbx	B	%r11
4	%rsp	C	%r12
5	%rbp	D	%r13
6	%rsi	E	%r14
7	%rdi	F	No register

Instruction Example: addq



- Add value in register rA to that in register rB
 - Store result in register rB
 - Note that Y86-64 only allows addition to be applied to register data
- Set condition codes based on result
- e.g., addq %rax,%rsi Encoding: 60 06
- Two-byte encoding
 - First indicates instruction type
 - Second gives source and destination registers


Arithmetic and Logical Operations

Instruction Code

Function Code

Add

`addq rA, rB`




Subtract (rA from rB)

`subq rA, rB`




And

`andq rA, rB`



Exclusive-Or

`xorq rA, rB`



- Refer to generically as "OPq"
- Encodings differ only by "function code"
 - Low-order 4 bits in the first instruction byte
- Set condition codes as side effect

Move Operations

Register → Register

`rrmovq rA, rB`



Immediate → Register

`irmovq V, rB`



Register → Memory

`rmmovq rA, D(rB)`



Memory → Register

`mrmovq D(rB), rA`



- Like the x86-64 `movq` instruction
- Simpler format for memory addresses
- Give different names to keep them distinct

Move Instruction Examples

X86-64

```
movq $0xabcd, %rdx
```

Encoding: 30 82 cd ab 00 00 00 00 00 00

```
movq %rsp, %rbx
```

Encoding: 20 43

```
movq -12(%rbp), %rcx
```

Encoding: 50 15 f4 ff ff ff ff ff ff

```
movq %rsi, 0x41c(%rsp)
```

Encoding: 40 64 1c 04 00 00 00 00 00 00

Y86-64

```
irmovq $0xabcd, %rdx
```

```
rrmovq %rsp, %rbx
```

```
mrmovq -12(%rbp), %rcx
```

```
rmmovq %rsi, 0x41c(%rsp)
```

Conditional Move Instructions

Move Unconditionally

`rrmovq rA, rB`




Move When Less or Equal

`cmovle rA, rB`




Move When Less

`cmovl rA, rB`




Move When Equal

`cmove rA, rB`



Move When Not Equal

`cmovne rA, rB`




Move When Greater or Equal

`cmovge rA, rB`



Move When Greater

`cmovg rA, rB`



- Refer to generically as "cmovXX"
- Encodings differ only by "function code"
- Based on values of condition codes
- Variants of `rrmovq` instruction
 - (Conditionally) copy value from source to destination register

Jump Instructions

Jump Unconditionally



Jump When Less or Equal



Jump When Less



Jump When Equal



Jump When Not Equal



Jump When Greater or Equal



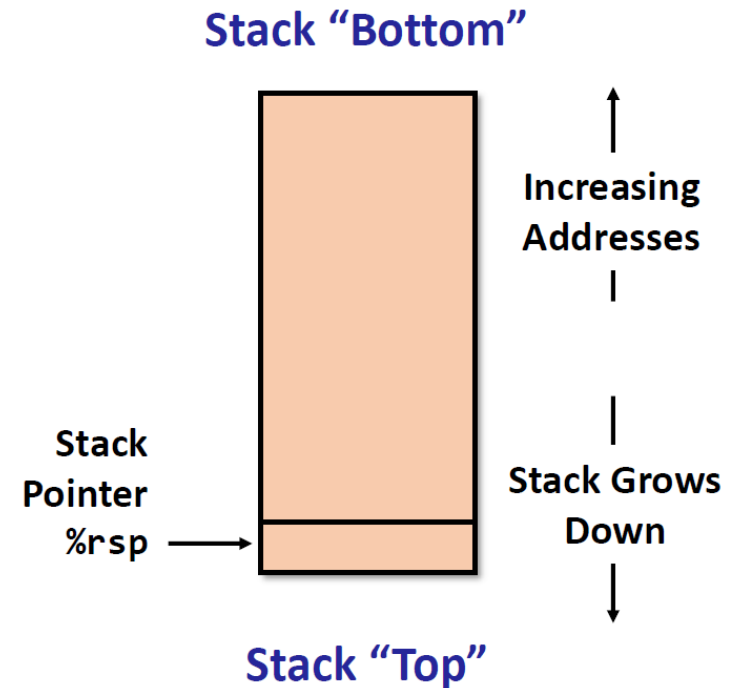
Jump When Greater



- Refer to generically as "jXX"
- Encodings differ only by "function code"
- Based on values of condition codes
- Same as x86-64 counterparts
- Encode full destination address
 - Unlike PC-relative addressing seen in x86-64

Y86 Program Stack

- Same as in x86-64
 - Region of memory holding program data
 - Used for supporting procedure calls
- Stack top indicated by %rsp
 - Address of top stack element
- Stack grows toward lower addresses
 - Top element is at lowest address in the stack
 - When pushing, must first decrement stack pointer
 - After popping, increment stack pointer



Stack Operations

- **pushq**

- Decrement %rsp by 8
- Store word from rA to memory at %rsp
- Like x86-64

pushq rA



- **popq**

- Read word from memory at %rsp
- Save in rA
- Increment %rsp by 8
- Like x86-64

popq rA



Subroutine Call and Return

- **call**

- Push address of next instruction onto stack
- Start executing instructions at Dest
- Like x86-64

call Dest

8	0	Dest
---	---	------

- **ret**

- Pop value from stack
- Use as address for next instruction
- Like x86-64

ret

9	0
---	---

Miscellaneous Instructions

- **nop**

- Don't do anything



- **halt**

- Stop executing instructions
- x86-64 has comparable instruction, but can't execute it in user mode
- We will use it to stop the simulator
- Encoding ensures that program hitting memory initialized to zero will halt



Status Conditions

- AOK
 - Normal operation
- HLT
 - Halt instruction encountered
- ADR
 - Bad address (either instruction or data) encountered
- INS
 - Invalid instruction encountered
- If AOK, keep going. Otherwise, stop program execution

Mnemonic	Code
AOK	1

Mnemonic	Code
HLT	2

Mnemonic	Code
ADR	3

Mnemonic	Code
INS	4

Outline

- Y86-64 ISA
 - Programmer-visible state
 - Instruction Format
 - Exceptions
- Y86-64 Programs
- CISC vs. RISC

Example: max.y

```
# Execution begins at address 0
.pos 0
irmovq stack, %rsp      # Set up stack pointer
call main               # Call main
halt                   # Terminate program

# Global data
.align 8
array:
    .quad 0x0000000000000003
    .quad 0x0000000000000004
    .quad 0x00000b000b000b00

main:
    irmovq array, %rbx    # %rbx <- array
    mrmovq (%rbx), %rdi    # %rdi <- array[0]
    mrmovq 8(%rbx), %rsi   # %rsi <- array[1]
    call max               # Call max
    rmmovq %rax, 16(%rbx)  # array[2] <- %rax
    ret

# long max(long x, y)
max:
    rrmovq %rdi, %rax      # %rax <- x
    subq %rsi, %rdi        # %rdi <- x - y; set flag
    cmovl %rsi, %rax       # if (x < y) %rax <- y
    ret

# The stack starts here and grows to lower addresses
.pos 0x200
stack:
```


Example: max.yo

- `yas max.ys`

```
0x000:                                     # Execution begins at address 0
0x000: 30f400020000000000000000          .pos 0
0x00a: 803000000000000000000000          irmovq stack, %rsp # Set up stack pointer
0x013: 00                                     call main # Call main
                                           halt # Terminate program


                                           # Global data
0x018:                                     .align 8
0x018:                                     array:
0x018: 030000000000000000000000          .quad 0x000000000000000003
0x020: 040000000000000000000000          .quad 0x000000000000000004
0x028: 000b000b000b000000000000          .quad 0x00000b000b000b00

0x030:                                     main:
0x030: 30f318000000000000000000          irmovq array, %rbx # %rbx <- array
0x03a: 507300000000000000000000          mrmovq (%rbx), %rdi # %rdi <- array[0]
0x044: 506308000000000000000000          mrmovq 8(%rbx), %rsi # %rsi <- array[1]
0x04e: 806200000000000000000000          call max # Call max
0x057: 400310000000000000000000          rmmovq %rax, 16(%rbx) # array[2] <- %rax
0x061: 90                                     ret

                                           # long max(long x, y)
0x062:                                     max:
0x062: 2070          rrmovq %rdi, %rax # %rax <- x
0x064: 6167          subq %rsi, %rdi # %rdi <- x - y; set flag
0x066: 2260          cmovl %rsi, %rax # if (x < y) %rax <- y
0x068: 90          ret

                                           # The stack starts here and grows to lower addresses
0x200:          .pos 0x200
0x200:          stack:
```

Example: Running max.yo on yis



```
$ yis max.yo
```

```
Stopped in 13 steps at PC = 0x13. Status 'HLT', CC Z=0 S=1  
O=0
```

```
Changes to registers:
```

```
%rax: 0x0000000000000000 0x0000000000000004
```

```
%rbx: 0x0000000000000000 0x0000000000000018
```

```
%rsp: 0x0000000000000000 0x0000000000000200
```

```
%rsi: 0x0000000000000000 0x0000000000000004
```

```
%rdi: 0x0000000000000000 0xffffffffffffffff
```

```
Changes to memory:
```

```
0x0028: 0x00000b000b000b00 0x0000000000000004
```

```
0x01f0: 0x0000000000000000 0x0000000000000057
```

```
0x01f8: 0x0000000000000000 0x0000000000000013
```

Sample C Program

```
1  long sum(long *start, long count)
2  {
3      long sum = 0;
4      while (count) {
5          sum += *start;
6          start++;
7          count--;
8      }
9      return sum;
10 }
```

x86-64 Code

```
long sum(long *start, long count)
start in %rdi, count in %rsi
1  sum:
2      movl    $0, %eax          sum = 0
3      jmp     .L2              Goto test
4  .L3:                          loop:
5      addq    (%rdi), %rax       Add *start to sum
6      addq    $8, %rdi          start++
7      subq    $1, %rsi          count--
8  .L2:                          test:
9      testq   %rsi, %rsi        Test sum
10     jne     .L3              If !=0, goto loop
11     rep; ret                  Return
```

Y86-64 Code

```
long sum(long *start, long count)
start in %rdi, count in %rsi
1  sum:
2      irmovq $8,%r8          Constant 8
3      irmovq $1,%r9          Constant 1
4      xorq %rax,%rax          sum = 0
5      andq %rsi,%rsi          Set CC
6      jmp     test            Goto test
7  loop:
8      mrmovq (%rdi),%r10      Get *start
9      addq %r10,%rax          Add to sum
10     addq %r8,%rdi            start++
11     subq %r9,%rsi            count--. Set CC
12  test:
13     jne     loop            Stop when 0
14     ret                     Return
```

Full Y86-64 Code

```
1  # Execution begins at address 0
2      .pos 0
3      irmovq stack, %rsp      # Set up stack pointer
4      call main               # Execute main program
5      halt                   # Terminate program
6
7  # Array of 4 elements
8      .align 8
9  array:
10     .quad 0x000d000d000d
11     .quad 0x00c000c000c0
12     .quad 0x0b000b000b00
13     .quad 0xa000a000a000
14
15  main:
16     irmovq array,%rdi
17     irmovq $4,%rsi
18     call sum                 # sum(array, 4)
19     ret
20
21  # long sum(long *start, long count)
22  # start in %rdi, count in %rsi
23  sum:
24     irmovq $8,%r8           # Constant 8
25     irmovq $1,%r9           # Constant 1
26     xorq %rax,%rax          # sum = 0
27     andq %rsi,%rsi          # Set CC
28     jmp     test            # Goto test
29  loop:
30     mrmovq (%rdi),%r10       # Get *start
31     addq %r10,%rax           # Add to sum
32     addq %r8,%rdi            # start++
33     subq %r9,%rsi            # count--. Set CC
34  test:
35     jne     loop             # Stop when 0
36     ret                     # Return
37
38  # Stack starts here and grows to lower addresses
39     .pos 0x200
40  stack:
```

Output of YAS Assembler

```
                                | # Execution begins at address 0
0x000:                          | .pos 0
0x000: 30f4000200000000000000 | irmovq stack, %rsp      # Set up stack pointer
0x00a: 8038000000000000000000 | call main                # Execute main program
0x013: 00                      | halt                     # Terminate program
                                |
                                | # Array of 4 elements
0x018:                          | .align 8
0x018:                          | array:
0x018: 0d000d000d000000      | .quad 0x000d000d000d
0x020: c000c000c0000000      | .quad 0x00c000c000c0
0x028: 000b000b000b0000      | .quad 0x0b000b000b00
0x030: 00a000a000a00000      | .quad 0xa000a000a000
                                |
0x038:                          | main:
0x038: 30f7180000000000000000 | irmovq array,%rdi
0x042: 30f6040000000000000000 | irmovq $4,%rsi
0x04c: 8056000000000000000000 | call sum                # sum(array, 4)
0x055: 90                    | ret
```

Output of YAS Assembler

```
                                | # long sum(long *start, long count)
                                | # start in %rdi, count in %rsi
0x056:                          | sum:
0x056: 30f808000000000000000000 |   irmovq $8,%r8          # Constant 8
0x060: 30f901000000000000000000 |   irmovq $1,%r9          # Constant 1
0x06a: 6300                      |   xorq %rax,%rax         # sum = 0
0x06c: 6266                      |   andq %rsi,%rsi         # Set CC
0x06e: 708700000000000000000000 |   jmp     test           # Goto test
0x077:                          | loop:
0x077: 50a700000000000000000000 |   mrmovq (%rdi),%r10      # Get *start
0x081: 60a0                      |   addq %r10,%rax         # Add to sum
0x083: 6087                      |   addq %r8,%rdi          # start++
0x085: 6196                      |   subq %r9,%rsi          # count--. Set CC
0x087:                          | test:
0x087: 747700000000000000000000 |   jne     loop           # Stop when 0
0x090: 90                       |   ret                   # Return
                                |
                                | # Stack starts here and grows to lower addresses
0x200:                          |   .pos 0x200
0x200:                          | stack:
```


Output of YIS Simulator

Stopped in 34 steps at PC = 0x13. Status 'HLT', CC Z=1 S=0 O=0

Changes to registers:

%rax:	0x0000000000000000	0x0000abcdabcdabcd
%rsp:	0x0000000000000000	0x0000000000000200
%rdi:	0x0000000000000000	0x0000000000000038
%r8:	0x0000000000000000	0x0000000000000008
%r9:	0x0000000000000000	0x0000000000000001
%r10:	0x0000000000000000	0x0000a000a000a000

Changes to memory:

0x01f0:	0x0000000000000000	0x0000000000000055
0x01f8:	0x0000000000000000	0x0000000000000013

Outline

- Y86-64 ISA
 - Programmer-visible state
 - Instruction Format
 - Exceptions
- Y86-64 Programs
- CISC vs. RISC

CISC (Complex Instruction Set Computer)

- Stack-oriented instruction set
 - Use stack to pass arguments, save program counter
 - Explicit push and pop instructions
- Arithmetic instructions can access memory
 - Requires memory read and write: e.g. `addq %rax, 12(%rbx,%rcx,8)`
 - Complex address calculation
- Condition codes
 - Set as side effect of arithmetic and logical instructions
- Philosophy
 - Add instructions to perform “typical” programming tasks
- DEC PDP-11 & VAX, IBM System/360, Motorola 68000, IA32, x86-64, ... (Dominant style through mid-80's)

RISC (Reduced Instruction Set Computer)

- Fewer, simpler instructions
 - Might take more to get given task done
 - Can execute them with small and fast hardware
 - Stanford MIPS, UCB RISC, Sun SPARC, IBM Power/PowerPC, ARM, SuperH, ...
- Register-oriented instruction set
 - Many more (typically 32+) registers
 - Use for arguments, return pointer, temporaries
- Only load and store instructions can access memory
 - Similar to Y86-64 `mrmovq` and `rmmovq`
- No condition codes
 - Test instructions return 0/1 in register

MIPS Registers

\$0	\$0	Constant 0	\$16	\$s0	
\$1	\$at	Reserved Temp.	\$17	\$s1	
\$2	\$v0	Return Values	\$18	\$s2	
\$3	\$v1		\$19	\$s3	
\$4	\$a0		\$20	\$s4	
\$5	\$a1	Procedure arguments	\$21	\$s5	
\$6	\$a2		\$22	\$s6	
\$7	\$a3		\$23	\$s7	
\$8	\$t0		\$24	\$t8	Caller Save Temp
\$9	\$t1		\$25	\$t9	
\$10	\$t2	Caller Save Temporaries: May be overwritten by called procedures	\$26	\$k0	Reserved for Operating Sys
\$11	\$t3		\$27	\$k1	
\$12	\$t4		\$28	\$gp	Global Pointer
\$13	\$t5		\$29	\$sp	Stack Pointer
\$14	\$t6		\$30	\$s8	Callee Save Temp
\$15	\$t7		\$31	\$ra	Return Address

MIPS Instruction Examples

R-R

Op	Ra	Rb	Rd	00000	Fn
----	----	----	----	-------	----

`addu $3,$2,$1` # Register add: $\$3 = \$2 + \$1$

Load/Store

Op	Ra	Rb	Offset
----	----	----	--------

`lw $3,16($2)` # Load Word: $\$3 = M[\$2+16]$

`sw $3,16($2)` # Store Word: $M[\$2+16] = \3

Branch

Op	Ra	Rb	Offset
----	----	----	--------

`beq $3,$2,dest` # Branch when $\$3 = \2

Jump

Op	Dest
----	------

`jmp Dest` # Jump to dest

CISC vs. RISC

- Original debate
 - CISC proponents – easy for compiler, fewer code bytes
 - RISC proponents – better for optimizing compilers, can make run fast with simple chip design
- Current status
 - For desktop processors, choice of ISA not a technical issue
 - With enough hardware, can make anything run fast
 - Code compatibility more important
 - x86-64 adopted many RISC features
 - More registers; use them for argument passing
 - For embedded processors, RISC makes sense
 - Smaller, cheaper, less power (e.g. most cell phones use ARM processor)

Summary

- Y86-64 Instruction Set Architecture
 - Similar state and instructions as x86-64
 - Simpler encodings
 - Somewhere between CISC and RISC
- How Important is ISA Design?
 - Less now than before
 - With enough hardware, can make almost anything go fast

Questions?