

Pipelining

'20H2


송 인 식

Outline

- Pipelining Concepts
- Pipeline Hazards


Sequential Processing

- Sequential processing: Wash-Dry-Fold-Store



Pipelined Processing

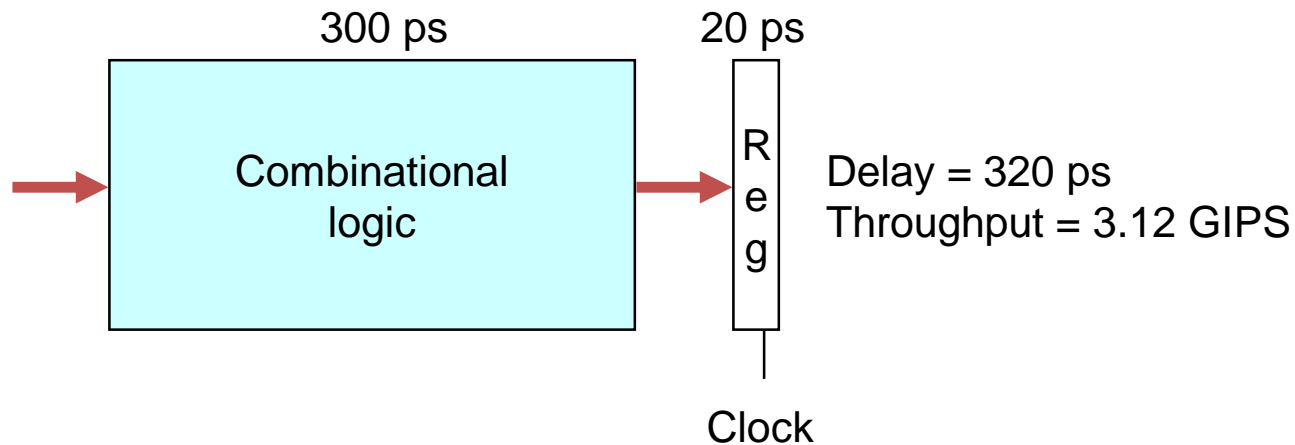
- Overlapping execution
- Parallelism improves performance



- Four loads:
 $\text{Speedup} = 8 / 3.5 = 2.3$
- Non-stop:
 $\text{Speedup} = 2n / (0.5n + 1.5) \approx 4$
= number of stages

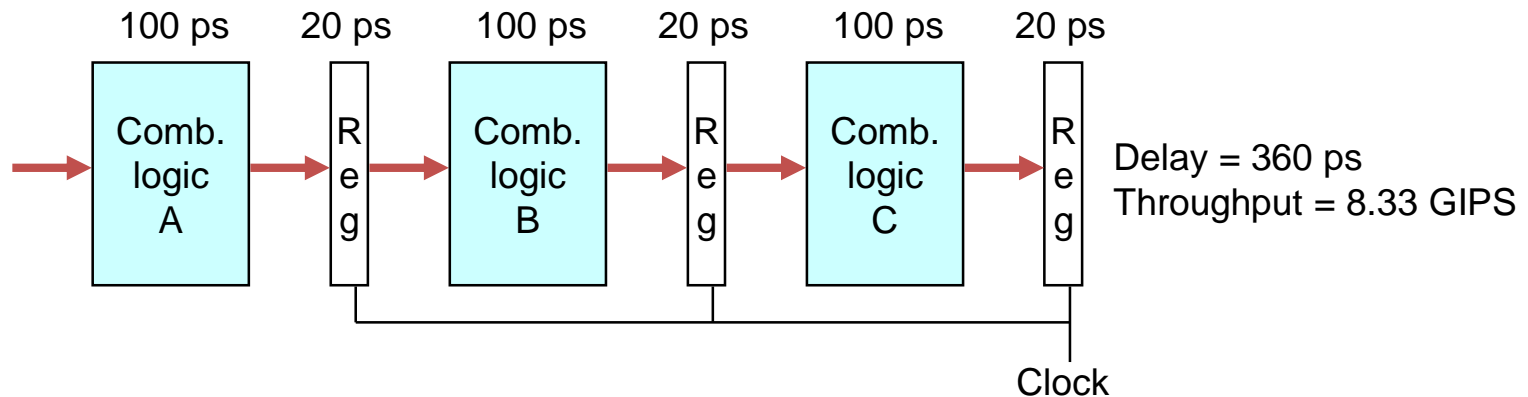
Sequential Execution

- An example CPU
 - Instruction execution requires total of 300 ps
 - Additional 20 ps to save results in registers
 - Must have clock cycle of at least 320 ps




Pipelined Execution

- Example: 3-way pipelined version
 - Divide combinational logic into 3 blocks of 100 ps each
 - Can begin new operation as soon as previous one passes through stage A
 - Begin new operation every 120 ps
 - Overall latency increases: 360 ps from start to finish
 - But, throughput improved by 2.67x




Pipeline Diagrams


- Unpipelined (= sequential)
 - Cannot start new operation until previous one completes



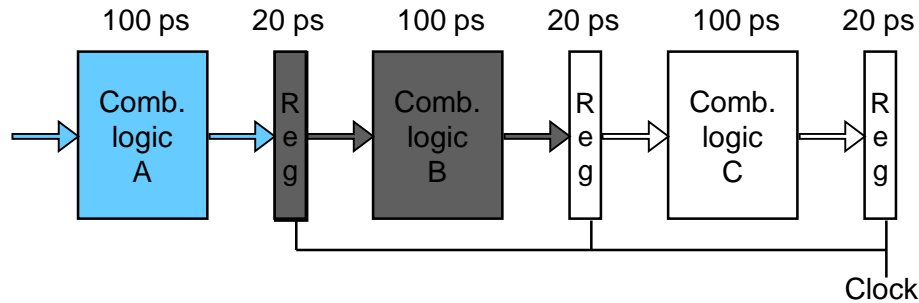
- 3-way pipelined
 - Up to 3 operations in process simultaneously



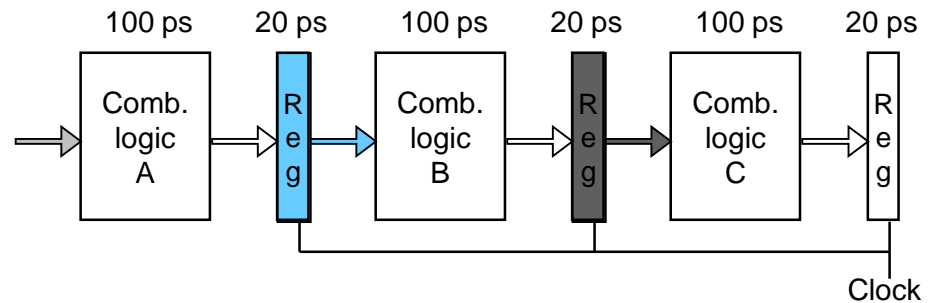
Operating a Pipeline




① Time = 239



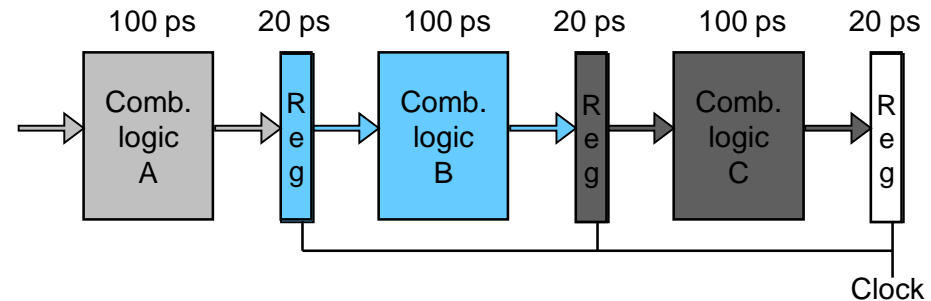
② Time = 241



③ Time = 300

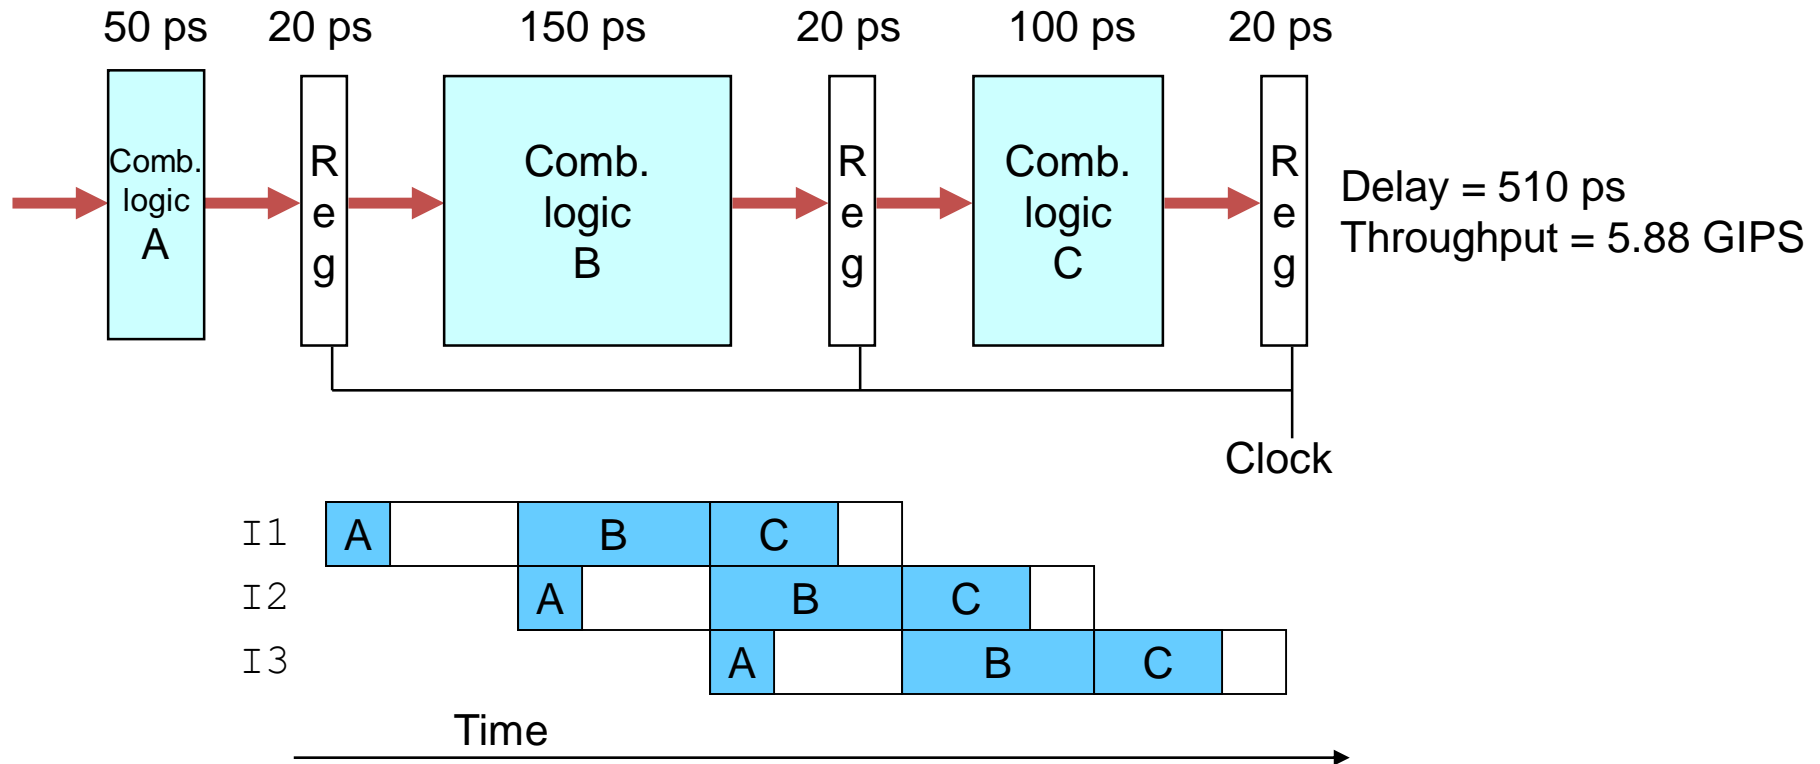


④ Time = 359



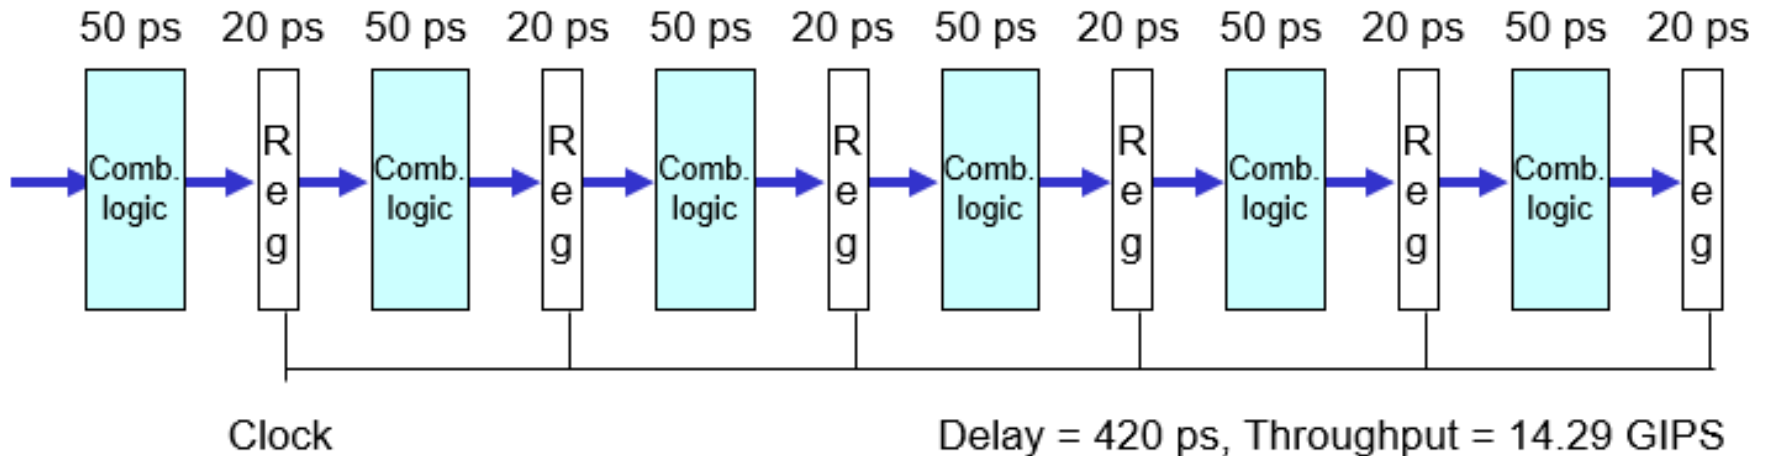
Pipelining

Limitations: Nonuniform Delays



- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages

Limitations: Register Overhead



- As pipeline deepens, overhead of loading registers becomes more significant
- Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%, 3-stage pipeline: 16.67%, 6-stage pipeline: 28.57%
- High speeds of modern processor designs obtained through very deep pipelining

Outline


- Pipelining Concepts
- Pipeline Hazards

Basic Steps of Execution


1. Instruction fetch step (F)
 2. Instruction decode/register fetch step (D)
 3. Execution/effective address step (E)
 4. Memory access (M)
 5. Register write-back step (W)
- Used in pipelined Y86-64 implementation later

Pipelined Instruction Execution

- Sequential Execution



- Pipelined Execution




```
addq %rcx, %rax  
subq %rdx, %rbx  
andq %rdx, %rcx
```

Major Hurdles (Hazards) of Pipelining

- Situations that prevent starting the next instruction in the next cycle
- **Structural** hazard
 - A required resource is busy
- **Data** hazard
 - Need to wait for previous instruction to complete its data read/write
- **Control** hazard
 - Deciding on control action depends on previous instruction

Structural Hazard




- Solution: resource duplication
 - Separate I and D caches for memory access conflict
 - Time-multiplexed or multi-port register file for register file access conflict

Data Hazard

- Read After Write (RAW) Hazard
 - An instruction tries to read operand before the previous instruction writes it
 - Called a “(true) dependence”
 - This hazard results from an actual need for communication

```
addq    %rdx, %rax
subq    %rax, %rcx
```



Solutions to Data Hazard


- Freezing the pipeline
- (Internal) Forwarding
- Compiler scheduling

Freezing The Pipeline

- Stall the pipeline until dependences are resolved
- ALU result to next instruction (3 stalls)

addq %rdx, %rax


subq %rax, %rcx



- Load result to next instruction (3 stalls)

mrmovq 0(%rdx), %rax

addq %rax, %rcx




(Internal) Forwarding

- Don't wait for the result to be stored in a register
- Requires extra connections in the datapath
- ALU result to next instruction (no stall)

addq %rdx, %rax

subq %rax, %rcx



- Load result to next instruction (1 stall)

mrmovq 0(%rdx), %rax


addq %rax, %rcx



Load interlock

Forwarding: More Example

```
addq    %rdx, %rax
subq    %rax, %rcx
andq    %rax, %rbx
addq    %rax, %rsi
xorq    %rax, %rdi
```




Compiler Scheduling

- Without compiler scheduling (1 stall)

`mrmovq 0(%rdx), %rax`

`addq %rax, %rcx`

`subq %rdi, %rsi`




- With compiler scheduling (no stall)

`mrmovq 0(%rdx), %rax`

`subq %rdi, %rsi`

`addq %rax, %rcx`



Control Hazard

- Caused by PC-changing instructions
 - (conditional/unconditional) Jump, Call / Return
- Stall-on-branch example: branch outcome determined in E stage
 - 2-cycle penalty for every branch
 - CPI (Cycles Per Instruction) = 1.3 for 15% branch frequency

Branch Instruction	F	D	E	M	W				
Branch successor	F	<i>stall</i>	F	D	E	M	W		
Branch successor + 1				F	D	E	M	W	
Branch successor + 2					F	D	E	M	W
Branch successor + 3						F	D	E	M
Branch successor + 4							F	D	E
Branch successor + 5								F	D

Solutions to Control Hazard


- Delayed branch
- Branch prediction

Delayed Branch


- Compiler approach
 - Branch delay slot filled by useful instructions
 - Branch delay slot becomes longer in modern processors
 - Microarchitecture-dependent

Delayed Branch


a. From before




Becomes




b. From target




Becomes



c. From fall-through



Becomes



Branch Prediction

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - (e.g., record recent history of each branch)
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history
 - Modern processors favor dynamic branch prediction

Branch Prediction Example

- Predict the branch outcome early in the pipeline
- Example: (static) predict-taken

Taken branch instruction	F	D	E	M	W			
Branch target	F	D	E	M	W			
Branch target + 1		F	D	E	M	W		
Branch target + 2			F	D	E	M	W	
Branch target + 3				F	D	E	M	W

Untaken branch instruction	F	D	E	M	W			
Branch target	F	D	idle	idle	idle			
Branch target + 1		F	idle	idle	idle	idle		
Untaken branch instruction + 1			F	D	E	M	W	
Untaken branch instruction + 2				F	D	E	M	W

Questions?