

Lecture Note 1: OS Introduction

February 26, 2021
Jongmoo Choi

Dept. of software
Dankook University

<http://embedded.dankook.ac.kr/~choijm>


(This slide is made by Jongmoo Choi. Please let him know when you want to distribute this slide)

J. Choi, DKU

Chap 1. A Dialog on the Book

■ OSTEP

- ✓ Operating Systems: Three Easy Pieces
- ✓ Homage to the Feynman's famous "Six Easy Pieces on Physics"
 - OS is about half as hard as Physics: from Six to Three Pieces



CONTENTS

Publisher's Note vii
Introduction by Paul Davies ix
Special Preface xix
Feynman's Preface xxv

ONE: Atoms in Motion 1
Introduction 1
Matter is made of atoms 4
Atomic processes 10
Chemical reactions 15

TWO: Basic Physics 23
Introduction 23
Physics before 1920 27
Quantum physics 33
Nuclei and particles 38

THREE: The Relation of Physics to Other Sciences 47

Introduction 47
Chemistry 48
Biology 49
Astronomy 50
Geology 61
Psychology 63
How did it get that way? 64

FOUR: Conservation of Energy 69
What is energy? 69
Gravitational potential energy 72
Kinetic energy 80
Other forms of energy 81

FIVE: The Theory of Gravitation 89
Planetary motions 89
Kepler's laws 90
Development of dynamics 92
Newton's law of gravitation 94
Universal gravitation 98
Cavendish's experiment 104
What is gravity? 107
Gravity and relativity 112

SIX: Quantum Behavior 115
Atomic mechanics 115
An experiment with bullets 117
An experiment with waves 120
An experiment with electrons 122
The interference of electron waves 124
Watching the electrons 127
First principles of quantum mechanics 133
The uncertainty principle 136

Index 139

Contents

- From Chap 1~2 of the OSTEP
- Chap 1. A Dialogue on the Book
- Chap 2. Introduction to Operating System
 - ✓ Virtualizing the CPU
 - ✓ Virtualizing Memory
 - ✓ Concurrency
 - ✓ Persistence
 - ✓ Design Goals
 - ✓ Some History
 - ✓ References

Chap 1. A Dialog on the Book

■ OSTEP

- ✓ What are Three Pieces: Virtualization, Concurrency, Persistence

Intro	Virtualization	Concurrency	Persistence	Appendices
Preface	1 Dialogue	12 Dialogue	25 Dialogue	35 Dialogue
TOC	4 Processes	13 Address Spaces	26 Concurrency and Threads	36 I/O Devices
1 Dialogue	5 Process API code	14 Memory API	27 Thread API	37 Hard Disk Drives
2 Introduction code	6 Direct Execution	15 Address Translation	28 Locks	38 Redundant Disk Arrays (RAID)
	7 CPU Scheduling	16 Segmentation	29 Locked Data Structures	39 Files and Directories
	8 Multi-level Feedback	17 Free Space Management	30 Condition Variables	40 File System Implementation
	9 Lottery Scheduling code	18 Introduction to Paging	31 Semaphores	41 Fast File System (FFS)
	10 Multi-CPU Scheduling	19 Translation Lookaside Buffers	32 Concurrency Bugs	42 FFSCK and Journaling
	11 Summary	20 Advanced Page Tables	33 Event-based Concurrency	43 Log-structured File System (LFS)
		21 Swapping Mechanisms	34 Summary	44 Flash-based SSDs
		22 Swapping Policies		45 Data Integrity and Protection
		23 Case Study: VAX/VMS		46 Summary
		24 Summary		47 Dialogue
				48 Distributed Systems
				49 Network File System (NFS)
				50 Andrew File System (AFS)
				51 Summary

(Source: <http://pages.cs.wisc.edu/~remzi/OSTEP/>)

Chap 1. A Dialog on the Book

■ OSTEP

✓ What to study?

Professor: They are the three key ideas we're going to learn about: virtualization, concurrency, and persistence. In learning about these ideas, we'll learn all about how an operating system works, including how it decides what program to run next on a CPU, how it handles memory overload in a virtual memory system, how virtual machine monitors work, how to manage information on disks, and even a little about how to build a distributed system that works when parts have failed. That sort of stuff.

Student: I have no idea what you're talking about, really.

Professor: Good! That means you are in the right class.

✓ How to study?

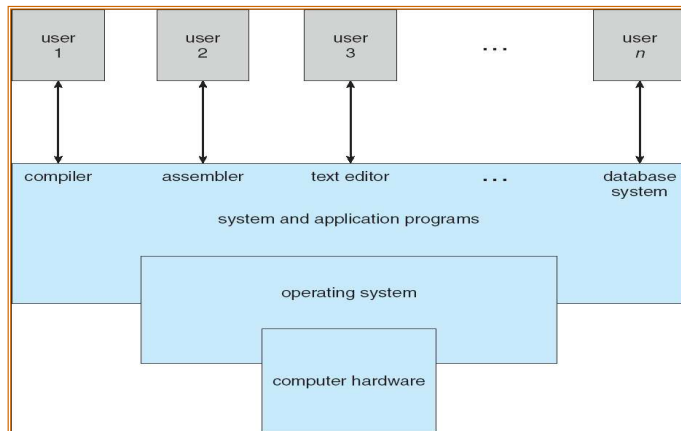
Student: I have another question: what's the best way to learn this stuff?

Professor: Excellent query! Well, each person needs to figure this out on their own, of course, but here is what I would do: go to class, to hear the professor introduce the material. Then, at the end of every week, read these notes, to help the ideas sink into your head a bit better. Of course, some time later (hint: before the exam!), read the notes again to firm up your knowledge. Of course, your professor will no doubt assign some homeworks and projects, so you should do those; in particular, doing projects where you write real code to solve real problems is the best way to put the ideas within these notes into action. As Confucius said...

Student: Oh, I know! 'I hear and I forget. I see and I remember. I do and I understand.' Or something like that.

Introduction

■ Layered structure of a computer system



(Source: A. Silberschatz, "Operating system Concept")

Chap 2. Introduction to Operating Systems

- 2.1 Virtualizing CPU
- 2.2 Virtualizing Memory
- 2.3 Concurrency
- 2.4 Persistence
- 2.5 Design Goals
- 2.6 Some history
- 2.7 Summary
- References

Introduction

■ What happens when a program runs?

✓ 1. Simple view about running a program




Figure 1.1 Computer Components: Top-Level View



(Source: W. Stalling, "Operating Systems: Internals and Design Principles")

Introduction

■ What happens when a program runs?

- ✓ Details: execute instructions

■ Fetch and Execute



Program Counter (PC) = Address of instruction
Instruction Register (IR) = Instruction being executed
Accumulator (AC) = Temporary storage

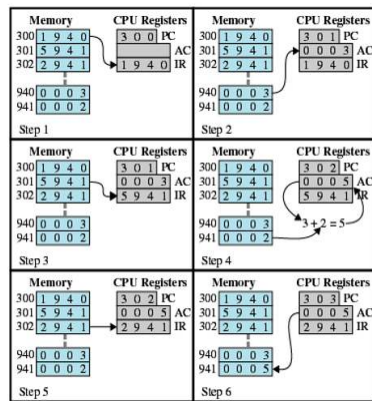
(c) Internal CPU registers

0001 = Load AC from Memory
0010 = Store AC to Memory
0101 = Add to AC from Memory

(d) Partial list of opcodes

<Hypothetical machine>

(Source: W. Stalling, "Operating Systems: Internals and Design Principles")



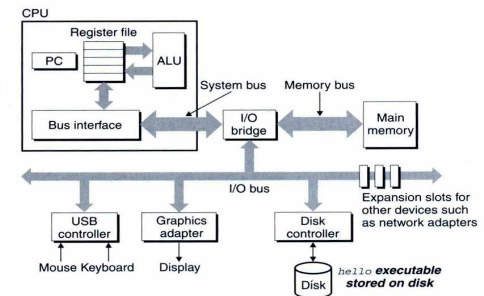
Introduction

■ What happens when a program runs?

- ✓ 2. A lot of stuff for running a program

- Loading, memory management, scheduling, context switching, I/O processing, file management, IPC, ...
- Operating system: 1) make it easy to run programs, 2) operate a system correctly and efficiently

Figure 1.4
Hardware organization
of a typical system.
CPU: Central
Processing Unit; ALU:
Arithmetic/Logic Unit; PC:
Program Counter; USB:
Universal Serial Bus.



(Source: computer systems: a programmer perspective)

Introduction

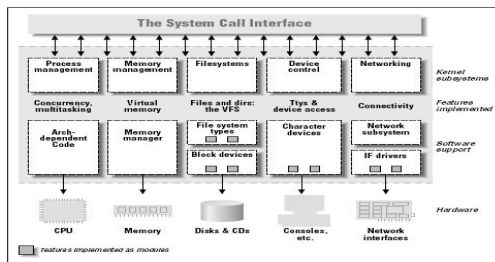
■ Definition of operating system

- ✓ Resource manager

- Physical resources: CPU (core), DRAM, Disk, Flash, KBD, Network, ...
- Virtual resources: Process, Thread, Virtual memory, Page, File, Directory, Driver, Protocol, Access control, Security, ...

- ✓ Virtualization (Abstraction)

- Transform a physical resource into a more general, powerful, and easy-to-use virtual form



(Source: Linux Device Driver, O'Reilly)

Introduction

■ System call

- ✓ Interfaces (APIs) provided by OS

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

(Source: A. Silberschatz, "Operating system Concept")

Introduction

■ System call

- ✓ Standard (e.g., POSIX, Win32, ...)
- ✓ **Mode switch** (user mode, kernel mode)

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

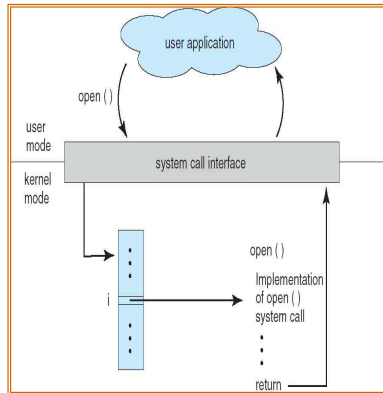
```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count)
```

return value	function name	parameters
<code>ssize_t</code>	<code>read</code>	<code>(int fd, void *buf, size_t count)</code>

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.



(Source: A. Silberschatz, "Operating system Concept")

2.1 Virtualizing CPU

■ A program for the discussion of virtualizing CPU

- ✓ call `Spin` (busy waiting and return when it has run for a second)
- ✓ print out a string passed in on the command line

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/time.h>
4 #include <assert.h>
5 #include "common.h"
6
7 int
8 main(int argc, char *argv[])
9 {
10     if (argc != 2) {
11         fprintf(stderr, "usage: cpu <string>\n");
12         exit(1);
13     }
14     char *str = argv[1];
15     while (1) {
16         Spin(1);
17         printf("%s\n", str);
18     }
19     return 0;
20 }
```

Figure 2.1: Simple Example: Code That Loops and Prints (`cpu.c`)

2.1 Virtualizing CPU

■ Execute the CPU program

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

■ Execute the program in parallel

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &
[1] 7353
[2] 7354
[3] 7355
[4] 7356
A
B
D
C
A
B
D
C
A
...
```

Process, Scheduling, ...

Figure 2.2: Running Many Programs At Once

2.1 Virtualizing CPU

■ Issues for Virtualizing CPU

- ✓ How to run a new program? → **process**
- ✓ How to make a new process? → `fork()`
- ✓ How to stop a process? → `exit()`
- ✓ How to execute a new process? → `exec()`
- ✓ How to block a process? → `sleep()`, `pause()`, `lock()`, ...
- ✓ How to select a process to run next? → **scheduling**
- ✓ How to run multiple processes? → context switch
- ✓ How to manage multiple cores (CPUs)? → multi-processor scheduling, cache affinity, load balancing
- ✓ How to communicate among processes? → IPC (Inter-Process Communication), socket
- ✓ How to notify an event to a process? → signal (e.g. `^C`)
- ✓ ...
- ✗ Illusion: A process has its own CPU even though there are less CPUs than processes



Quiz for 1th-Week 2st-Lesson

Quiz

- ✓ 1. Operating system is defined as a resource manager. What kinds of resources are managed by operating system? Discuss physical and virtual resources separately.
- ✓ 2. What is the role of "&" in the below example? (I do this experiment using wsl(windows subsystem for Linux) in my laptop.)
- ✓ Due: until 6 PM Friday of this week (5th, March)

```

choijm@DESKTOP-7SHQTVH: ~/OS_exam
choijm@DESKTOP-7SHQTVH:~/OS_exam$ ls
common.h  cpu.c
choijm@DESKTOP-7SHQTVH:~/OS_exam$ gcc -o cpu cpu.c
choijm@DESKTOP-7SHQTVH:~/OS_exam$ ./cpu
Usage: cpu <string>
choijm@DESKTOP-7SHQTVH:~/OS_exam$ ./cpu A
A
choijm@DESKTOP-7SHQTVH:~/OS_exam$ ./cpu A & ./cpu B & ./cpu C &
[1] 4459
[2] 4460
[3] 4461
choijm@DESKTOP-7SHQTVH:~/OS_exam$ B
A
B
C

```

2.2 Virtualizing Memory

Execute the Mem program

```

prompt> ./mem
(2134) address pointed to by p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C

```

Execute the program in parallel

```

prompt> ./mem & ./mem &
[1] 24113
[2] 24114
(24113) address pointed to by p: 0x200000
(24114) address pointed to by p: 0x200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
(24113) p: 4
(24114) p: 4
...

```

Figure 2.4: Running The Memory Program Multiple Times

Same address but independent

2.2 Virtualizing Memory

Memory

- ✓ Can be considered as an array of bytes

Another program example

- ✓ Allocate a portion of memory and access it

```

1 #include <unistd.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include "common.h"
5
6 int
7 main(int argc, char *argv[])
8 {
9     int *p = malloc(sizeof(int));           // a1
10    assert(p != NULL);
11    printf("(d) address pointed to by p: %p\n",
12           getpid(), p);                     // a2
13    *p = 0;                                  // a3
14    while (1) {
15        Spin(1);
16        *p = *p + 1;
17        printf("(d) p: %d\n", getpid(), *p); // a4
18    }
19    return 0;
20 }

```

Figure 2.3: A Program That Accesses Memory (mem.c)

2.2 Virtualizing Memory

Issues for Virtualizing Memory

- ✓ How to manage the address space of a process? → text, data, stack, heap, ...
- ✓ How to allocate memory to a process? → malloc(), calloc(), brk(), ...
- ✓ How to deallocate memory from a process? → free()
- ✓ How to manage free space? → buddy, slab, ...
- ✓ How to protect memory among processes? → virtual memory
- ✓ How to implement virtual memory? → page, segment
- ✓ How to reduce the overhead of virtual memory? → TLB
- ✓ How to share memory among processes? → shared memory
- ✓ How to exploit memory to hide the storage latency? → page cache, buffer cache, ...
- ✓ How to manage NUMA? → local/remote memory
- ✓ ...

Illusion: A process has its own unlimited and independent memory even though several processes are sharing limited memory in reality

2.3 Concurrency

- Background: how to create a new scheduling entity?
 - ✓ Two programming model: process (task) and thread
 - ✓ Key difference: data sharing

```
// fork example (Refer to the Chapter 5 in OSTEP)
// by J. Choi (choijm@dku.edu)
#include <stdio.h>
#include <stdlib.h>

int a = 10;

void *func()
{
    a++;
    printf("pid = %d\n", getpid());
}

int main()
{
    int pid;
    if ((pid = fork()) == 0) { //need exception handle
        func();
        exit(0);
    }
    wait();
    printf("a = %d by pid = %d\n", a, getpid());
}
```

(Source: System programming lecture note)

```
// thread example (Refer to the Chapter 27 in OSTEP)
// by J. Choi (choijm@dku.edu)
#include <stdio.h>
#include <stdlib.h>

int a = 10;

void *func()
{
    a++;
    printf("pid = %d\n", getpid());
}

int main()
{
    pthread_t p_thread;
    if ((pthread_create(&p_thread, NULL, func, (void *)NULL))
        < 0) {
        exit(0);
    }
    pthread_join(p_thread, (void *)NULL);
    printf("a = %d by pid = %d\n", a, getpid());
}
```

J. Choi, DKU

2.3 Concurrency

- Concurrency
 - ✓ Problems arise when working on many things simultaneously on the same data
- A program for discussing concurrency

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include "common.h"
4
5 volatile int counter = 0;
6 int loops;
7
8 void *worker(void *arg) {
9     int i;
10    for (i = 0; i < loops; i++) {
11        counter++;
12    }
13    return NULL;
14 }
15
16 int
17 main(int argc, char *argv[])
18 {
19     if (argc != 2) {
20         fprintf(stderr, "usage: threads <value>\n");
21         exit(1);
22     }
23     loops = atoi(argv[1]);
24     pthread_t p1, p2;
25     printf("Initial value : %d\n", counter);
26
27     pthread_create(&p1, NULL, worker, NULL);
28     pthread_create(&p2, NULL, worker, NULL);
29     pthread_join(p1, NULL);
30     pthread_join(p2, NULL);
31     printf("Final value : %d\n", counter);
32     return 0;
33 }
```

Figure 2.5: A Multi-threaded Program (threads.c)



22

J. Choi, DKU

2.3 Concurrency

- Execute the multi-thread program

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value   : 2000
```

```
prompt> ./thread 100000
Initial value : 0
Final value   : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value   : 137298 // what the??
```

- ✓ Programing model
 - thread model: share data section (a.k.a data segment)
 - process model: independent, need explicit IPC for sharing
- ✓ Reason for the odd results for the large loop
 - Lack of **atomicity**, scheduling effect, ... → need concurrency control



23

J. Choi, DKU

2.3 Concurrency

- Issues for Concurrency
 - ✓ How to support concurrency correctly? → **lock()**, **semaphore()**
 - ✓ How to implement atomicity in hardware? → **test_and_set()**, **swap()**
 - ✓ What is the semaphore?
 - ✓ What is the monitor?
 - ✓ How to solve the traditional concurrent problems such as **producer-consumer**, **readers-writers** and **dining philosophers**?
 - ✓ What is a deadlock?
 - ✓ How to deal with the **deadlock**?
 - ✓ How to handle the **timing bug**?
 - ✓ What is the asynchronous I/Os?
 - ✓ ...
- ➡ Illusion: Multiple processes run in a cooperative manner on shared resources even though they actually race with each other on the resources




24


J. Choi, DKU

2.4 Persistence

■ Background: DRAM vs. Disk



- ✓ Capacity, Speed, Cost, ...
- ✓ Access granularity: Byte vs. Sector
- ✓ Durability: Volatile vs. Non-volatile



2.4 Persistence

■ Persistence

- ✓ Users want to maintain data permanently (durability)
- ✓ DRAM is volatile, requiring write data into storage (disk, SSD) **explicitly**

■ A program for discussing persistence

- ✓ Use the notion of a file (not handle disk directly)

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <assert.h>
4 #include <fcntl.h>
5 #include <sys/types.h>
6
7 int
8 main(int argc, char *argv[])
9 {
10     int fd = open("/tmp/file", O_WRONLY | O_CREAT | O_TRUNC, S_IRWXU);
11     assert(fd > -1);
12     int rc = write(fd, "hello world\n", 13);
13     assert(rc == 13);
14     close(fd);
15     return 0;
16 }
```

Figure 2.6: A Program That Does I/O (io.c)

2.4 Persistence

■ Issues for Persistence

- ✓ How to access a file? → `open()`, `read()`, `write()`, ...
- ✓ How to manage a file? → **inode**, FAT, ...
- ✓ How to manipulate a directory?
- ✓ How to design a file system? → UFS, LFS, **Ext2/3/4**, FAT, F2FS, NFS, AFS, ...
- ✓ How to find a data in a disk?
- ✓ How to improve performance in a file system? → cache, delayed write, ...
- ✓ How to handle a fault in a file system? → **journaling**, **copy-on-write**
- ✓ What is a role of a disk device **driver**?
- ✓ What are the internals of a disk and SSD?
- ✓ What is the RAID?

➤ Illusion: Data is always maintained in a reliable non-volatile area while it is often kept in a volatile DRAM (for performance reason) and storage is broken from time to time.

2.5 Design Goals

■ Abstraction

- ✓ Focusing on relevant issues only while hiding details
 - E.g. Car, File system, Make a program without thinking of logic gates
- ✓ "Abstraction is fundamental to everything we do in computer science" by Remzi

■ Performance

- ✓ Minimize the overhead of the OS (both time and space)

■ Protection

- ✓ Isolate processes from one another
- ✓ Access control, security, ...

■ Reliability

- ✓ Fault-tolerant


■ Others

- ✓ Depend on the area where OS is employed
- ✓ Real time, Energy-efficiency, Mobility, Load balancing, Autonomous, ...

2.5 Design Goals

- Separation of Policy and Mechanism

- ✓ Policy: Which (or What) to do?
 - e.g.) Which process should run next?
- ✓ Mechanism: How to do?
 - e.g.) Multiple processes are managed by a scheduling queue or RB-tree



(Source: Security Principles and Policies CS 236 On-Line MS Program Networks and Systems Security, Peter Reiher, Spring, 2008)

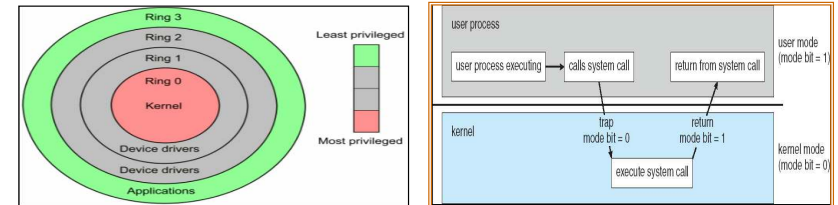
2.6 Some History

- Early Operating Systems: Just libraries

- ✓ Commonly-used functions such as low-level I/Os (e.g. MS-DOS)
- ✓ **Batch processing**
 - a number of jobs were set up and then run all together (Not interactive)

■ Beyond Libraries: Protection

- ✓ Require OS to be treated differently than user applications
- ✓ Separation user/kernel mode, system call
- ✓ Use **trap** (special instruction, SW interrupt) to go into the kernel mode
 - Transfer control to a pre-specific trap handler (system_call handler)




(Source: Google Image)

(Source: A. Silberschatz, "Operating system Concept")

2.6 Some History

- The Era of Multiprogramming (c.f. multitasking)

- ✓ Definition: OS load a number of applications into memory and switch them rapidly
- ✓ Reason: Advanced hardware → Want to utilize machine resources better → Multiple users share a system (workstation, minicomputer) → multiprogramming (and multitasking)
- ✓ Especially important due to the slow I/O devices → while doing I/O, switch CPU to another process → enhancing CPU utilization
- ✓ Memory protection and concurrency become quite important → UNIX

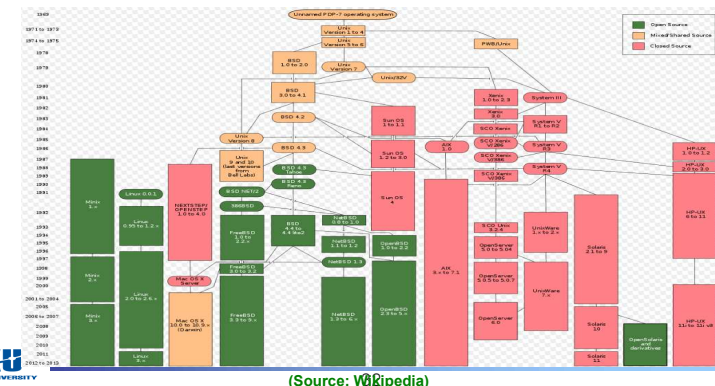


(Source: Google Image)

2.6 Some History

- The Era of Multiprogramming (c.f. multitasking)

- ✓ UNIX
 - By Ken Thompson and Dennis Ritchie (Bell Labs), Influenced by Multics
 - C language based, excellent features such as shell, pipe, inode, small, everything is a file, ...
 - Influence OSes such as BSD, SUNOS, AIX, HPUX, Nextstep and Linux



(Source: Wikipedia)

2.6 Some History

■ The Modern Era

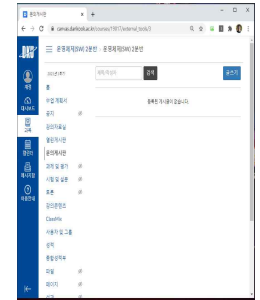
- ✓ PC
 - MS Windows, Mac OS X, Linux, ...
- ✓ Smartphone
 - Android, iOS, Windows Mobile, ...
- ✓ IoT
 - What is the next?



2.7 Summary

■ OS

- ✓ Resource manager (Efficiency)
- ✓ Make systems easy to use (Convenience)
- Cover in this book
 - ✓ Virtualization, Concurrency, Persistence
- Not being covered
 - ✓ Network, Security, Graphics
 - ✓ There are several excellent courses for them



- Homework 1: summarize the chap 2 of the OSTEP
 - Requirement: 1) personal, 2) up to 6 pages for summary, 3) 1 page for the goal you want to study
 - Due: until 6 PM, 19th March (Friday)
 - Bonus: Snapshot of the results of example programs in a Linux system (ubuntu on virtual box or wsl or Linux server)
- Any questions? Feel free to put your questions at "문의게시판"!!



Quiz for 2th-Week 1st-Lesson

■ Quiz


- ✓ 1. What are the differences between disks and DRAM? (at least 3 differences). These differences lead operating system to manage them differently (memory object vs. file)
- ✓ 2. Discuss differences between interrupt and trap which was discussed in page 30.
- ✓ Due: until 6 PM Friday of this week (12th, March)

TRAP	INTERRUPT
A signal raised from a user program that indicates the operating system to perform on some functionality immediately	A signal to the processor emitted by hardware indicating an event that needs immediate attention
Generated by an instruction in the user program	Generated by hardware devices
Invokes OS functionality - it transfers the control to the trap handler	Triggers the processor to execute the corresponding interrupt handler routine
Synchronous and can arrive after the execution of any instruction	Asynchronous and can occur at the execution of any instruction
Also called a software interrupt	Also called a hardware interrupt


Visit www.PEDIAA.com

Appendix


■ OS structure in General




(Source: Operating System Concepts)



(Source: Linux Device Driver)



(Source: <https://www.cs.rutgers.edu/~pxk/416/notes/03-concepts.html>)



(Source: Modern Operating System)